PRODUCT REQUIREMENTS DOCUMENT

Distill

The AI Agent That Turns Your Chaos Into Your Voice in 3 Minutes

Version	1.0 — Draft
Date	March 2026
Status	For Review
Product Type	Agentic Creator Tool — Web App
Primary User	General creators, vloggers, content entrepreneurs
Core Output	3-minute, creator-voiced content brief from any raw input

1. Product Overview

Distill is an AI agent that does two things most tools cannot do together: it learns how you talk, and it compresses what you think.

Every creator has the same problem. The raw material — the voice note at 2am, the scattered Google Doc, the 45-minute rant that had three brilliant points buried inside it — never makes it to the camera clean. The ideas are real. The voice is real. But the gap between unfiltered thought and a tight, deliverable 3-minute piece costs creators hours they don't have and a lot of good content that never gets made.

Distill closes that gap with four agents working in sequence. The first two understand your input — whatever format it comes in. The third extracts and compresses the ideas into a 3-minute structure. The fourth shapes the output in your specific voice, learned from your own content and the creators you admire. The result is a ready-to-deliver script or structured brief that sounds exactly like you — because it is built from you.

The Core Insight

Other AI writing tools make content that sounds like AI. Distill makes content that sounds like you. The Voice Profile is not a prompt — it is a persistent, ever-improving model of the user's specific tone, rhythm, vocabulary, energy, and style, trained on their own videos and the reference creators they feed it. Every output is calibrated to that profile. The longer you use Distill, the more it sounds like you.

2. Problem Statement

2.1 The Creator's Friction Loop

The modern creator's workflow has a structural problem that no tool has fully solved:

- Raw ideas come out messy — voice memos, scattered notes, video rambles, half-finished scripts
- Editing those ideas into deliverable content takes hours of cognitive work that isn't 'creating'
- AI writing tools help with structure but produce generic output that sounds nothing like the creator's voice
- The creator then spends more time editing the AI output to sound like themselves than they would have spent writing from scratch
- The 3-minute format is particularly brutal — tight enough that every word must earn its place, long enough that structure matters

The result: creators who have more to say than they can deliver. Good ideas die in the notes app. Content calendars slip. The creator shows up to film without a clear through-line and wings it — which wastes production time and often produces content they're not proud of.

2.2 Why Existing Tools Miss

ChatGPT / Claude (general LLMs)	Can compress text, can help structure	No persistent voice model. Output sounds like AI unless the user writes a very detailed style prompt every time. Cannot ingest video or audio natively.
Descript / Otter.ai	Transcribes audio and video accurately	Does nothing with the transcript — no compression, no structure, no voice modeling. Output is a wall of raw text.
Notion AI / Jasper	Can restructure written notes, create outlines	No voice learning, no audio/video ingestion, no 3-minute timing constraint. Generic editorial voice by default.
Teleprompter apps	Helps deliver a script	Assumes you already have a tight script. Does nothing to help create it.
Human editors	Understand creator voice deeply	Expensive, slow, not available at 2am when the idea hits. Doesn't scale.

3. User Personas

The Idea-Rich Creator	Prolific thinker. Always has something to say. Has 200 voice memos and 47 unfinished Google Docs. Posts inconsistently because prep time is the bottleneck.	Voice memos, rambling video drafts, scattered bullet points	Something that takes the chaos and gives them a clean, structured 3-minute brief they can film immediately
The Consistency Seeker	Posts regularly but feels like their content sounds 'off' — like they're reading AI output instead of talking. Has tried ChatGPT but hates how it sounds.	Scripts they wrote themselves or lightly edited AI output	A tool that makes the polished version sound as natural and specific as their unpolished version
The Multi-Platform Creator	Repurposes content across YouTube, TikTok, Instagram. Needs a 3-minute version, a 60-second version, and a written version of the same idea.	Longer-form scripts, full YouTube videos they want to cut down	Fast compression into the 3-minute tier with their voice intact
The New Creator	Just starting. Knows their ideas but doesn't have a 'voice' yet. Feeds in reference creators they want to sound like alongside their own raw material.	Mostly notes and brainstorms, minimal own video content yet	Reference-based voice building that helps them find and refine their style faster

4. The Voice Profile — Core Concept

4.1 What It Is

The Voice Profile is the heart of Distill. It is a persistent, structured model of the user's unique creative identity — built by analyzing their own videos and reference content they admire, and refined with every piece of output the user accepts, edits, or rejects.

It is not a saved prompt. It is not a style guide the user writes. It is an inferred model — assembled by the agent by watching and reading how the creator actually communicates, not how they think they communicate.

4.2 Voice Attributes Extracted

Vocabulary level	Simple / Conversational / Technical / Mixed
Sentence rhythm	Short punchy bursts / Long flowing sentences / Mixed cadence
Energy signature	High energy / Calm and measured / Builds to peaks / Deadpan
Opener pattern	Direct hook / Story / Question / Statement of fact / Vibe-set
Transition style	Hard cuts / Smooth bridges / Rhetorical questions / Callback loops
Closer pattern	Call to action / Reflection / Punchline / Mic-drop statement
Filler tendencies	Specific phrases, verbal tics, breath patterns, pause placement
Cultural voice markers	Slang, references, humor type, regional flavor
Pacing (words/min)	Slow (<100) / Medium (100-140) / Fast (140-180) / Variable
Authenticity markers	Personal pronouns, vulnerability moments, direct address ('you', 'we')

4.3 How the Profile Is Built

- User uploads their own existing videos (YouTube links, direct uploads, TikTok links) — minimum 3 recommended for initial calibration
- User uploads or links reference creators they admire — 'I want to capture this person's energy' or 'This is the tone I'm going for'
- Voice Ingestion Agent transcribes all content, parses it into the attribute matrix above, and builds the initial Voice Profile
- The profile is not static — every time the user accepts an output, edits it significantly, or rates a draft, the feedback loop updates the profile weights
- User can view their Voice Profile in plain language: 'Your content tends to open with a direct statement, builds through personal story, and closes with a reframe. Your average sentence is 11 words. You tend to address your audience directly about 4x per minute.'
- User can manually adjust any attribute — 'I want to be more conversational' or 'Less slang in this series' — and the override is respected for all future outputs

4.4 Reference Content vs. Own Content

The platform distinguishes between two input types in the Voice Profile:

- Own Content — the creator's own published videos, recordings, and writings. These carry full weight in the profile. The agent learns from what the creator actually does.
- Reference Content — creators the user admires, videos with the energy or tone they're going for. These carry partial weight — used as directional input, not as a source of vocabulary or identity. The agent absorbs the energy and structural patterns of references, not their words.

Why the Reference Layer Matters

New creators don't have enough of their own content to calibrate a profile yet. The reference layer gives them a starting point. A creator might say 'I want the directness of Alex Hormozi, the humor of Quinta Brunson, and the cultural fluency of Druski.' Distill doesn't copy any of them — it extracts the structural and tonal patterns from each reference and blends them into a target profile the creator can evolve toward.

5. Agentic Architecture

5.1 Overview

Distill runs four agents in sequence for every output. Agents 1 and 2 run in parallel when multiple input types are submitted simultaneously. Agent 3 is the creative core — it holds the 3-minute constraint as a hard rule and structures accordingly. Agent 4 is the voice layer — it never changes the substance of what Agent 3 built, only how it sounds.

The user's Voice Profile is injected into every call to Agent 3 and Agent 4 as a structured context object, not a free-text prompt. This ensures consistency across sessions and prevents profile drift over time.

Pipeline: [Input Ingestion Agent] + [Reference Ingestion Agent] → [Idea Extraction Agent] → [Compression & Structure Agent] → [Voice Calibration Agent] → Output

5.2 Agent Specifications

Agent 01 — Convert any raw input format into structured idea material

Trigger: User submits new raw content: writes text, uploads audio/voice memo, uploads or links a video, or submits bullet points

Inputs: Raw input in any format: typed text (script, notes, brainstorm), audio file (mp3, m4a, wav), video file (mp4, mov), or YouTube/TikTok URL

Outputs: Structured Idea Dump: a cleaned, timestamped transcription (for audio/video), identified distinct idea threads, emotional tone markers per section, rough word count and estimated duration of raw material, confidence score on transcription quality

Tools / APIs: Claude API (text understanding, idea thread identification), Whisper API / AssemblyAI (audio and video transcription), yt-dlp (YouTube URL extraction), TikTok scraper (TikTok URL extraction), FFmpeg (audio extraction from video uploads)

Priority: P0 — Core. Foundation of the entire pipeline. Must handle all 5 input types reliably before other agents run.

Agent 02 — Analyze creator content and reference media to build and update the Voice Profile

Trigger: User uploads their own past content (initial setup) OR reference creator content (initial setup) OR user accepts/rejects an output draft (ongoing learning signal)

Inputs: Creator's own videos and recordings (for own-content analysis), reference creator videos and content (for reference analysis), user feedback on past outputs (accept / edit / reject + what changed)

Outputs: Updated Voice Profile object: all 10 voice attributes scored and described in plain language, weighted blend of own content vs. reference content, version history with change log, plain-language profile summary the user can read and edit

Tools / APIs: Claude API (tone analysis, style extraction, attribute scoring), Whisper API (transcription of ingested videos), YouTube Data API (metadata + captions for reference creators), TikTok scraper, internal Voice Profile data model (Postgres)

Priority: P0 — Core for initial setup. Runs in background continuously as user submits feedback.

Agent 03 — Extract the 3-4 central ideas and build a tight 3-minute structure

Trigger: Structured Idea Dump from Agent 01; user's Voice Profile from Agent 02; optional user intent signal ('this is for YouTube intro' / 'this is a standalone piece' / 'this is a response to a trend')

Inputs: Full Idea Dump, Voice Profile object, user's stated intent (if provided), word-per-minute pacing value from Voice Profile

Outputs: 3-Minute Content Blueprint: a hook (opening 15-20 seconds), 2-3 core idea blocks with a connecting thread, a closer (final 20-30 seconds), total word count calibrated to user's pacing (e.g., at 130 wpm = ~390 words), estimated runtime per section, a 'what got cut and why' note (so the user understands what the agent deprioritized and can override), alternate opener options (2 variants)

Tools / APIs: Claude API (compression reasoning, structural editing, priority ranking of ideas), Voice Profile pacing attribute (timing calibration), 3-minute constraint enforced as a hard rule — output never exceeds 410 words or falls below 340 words at standard pacing

Priority: P0 — Core. The 3-minute constraint is non-negotiable. The agent must justify every inclusion and every cut.

Agent 04 — Rewrite the structure in the creator's specific voice

Trigger: 3-Minute Content Blueprint from Agent 03; full Voice Profile object; user's pacing, vocabulary level, energy signature, opener pattern, closer pattern, and filler tendencies

Inputs: Content Blueprint (structure and ideas), Voice Profile (all 10 attributes), any user overrides applied this session

Outputs: Final Distilled Script: complete 3-minute script written in the creator's voice with: opening in the creator's typical hook style, body in their natural sentence rhythm and vocabulary, closer in their typical close pattern, inline timing markers ([0:00] [0:30] [1:00] etc.), a readability score (how conversational vs. formal), a voice match confidence score ('This output is 87% aligned with your Voice Profile'), diff view vs. a generic version so user can see what the voice layer changed

Tools / APIs: Claude API (voice-aware rewriting, style transfer from profile, conversational naturalness), Voice Profile object (injected as structured context, not free-text prompt), user override flags

Priority: P0 — Core. This is the most differentiated step in the pipeline. Output must not sound like AI. If voice match confidence falls below 70%, agent flags the output and asks user to review Voice Profile calibration.

6. Feature Requirements

6.1 Input Handling

- System MUST accept all five input types in a single session: typed text, audio upload, video upload, YouTube URL, TikTok URL
- System MUST handle mixed inputs — user can submit a voice memo AND a bullet list in the same session and the agent merges them

- System MUST surface a transcription preview before running compression — user can correct transcription errors before the pipeline proceeds
- System MUST detect and flag very low-quality audio (background noise, unclear speech) and warn the user before running transcription
- System SHOULD support drag-and-drop file upload and mobile share-sheet (sharing a voice memo from iPhone to Distill directly)
- System MUST support YouTube and TikTok URL paste — no download required by the user

6.2 Voice Profile Setup

- System MUST include a Voice Profile onboarding wizard: step-by-step guide to uploading own content and reference content
- System MUST require a minimum of 3 own-content videos OR 15 minutes of own audio for initial calibration — below this threshold, the profile is marked as 'low confidence'
- System MUST allow reference content tagging: user marks what they want to borrow from each reference ('energy,' 'structure,' 'humor,' 'directness')
- System MUST display the Voice Profile in plain language after calibration: 'Here is what we learned about how you communicate — review and adjust anything that feels off'
- System MUST allow manual attribute overrides: user can edit any Voice Profile attribute directly
- System MUST support multiple Voice Profiles for creators with distinct content series (e.g., 'Casual Vlog Voice' vs. 'Educational Series Voice')
- System SHOULD surface a Voice Profile update notification when enough new feedback has accumulated to meaningfully change the profile: 'We have enough new data to improve your voice model — want to update?'

6.3 The 3-Minute Output

- System MUST enforce a hard word count ceiling and floor based on user's pacing attribute — default calibrated to 130 wpm standard conversational pace
- System MUST display inline timing markers throughout the output: [0:00] HOOK [0:20] POINT 1 [1:10] POINT 2 etc.
- System MUST surface a 'What Got Cut' panel explaining which ideas were deprioritized and why — user can reinstate any cut idea
- System MUST offer 2 alternate opener variants in every output — user can swap to either or request a third
- System MUST display voice match confidence score and which attributes drove or hurt the score
- System MUST show a diff between the final output and what a generic (non-voice-calibrated) version would look like — so users understand the value of their Voice Profile
- System SHOULD offer a 'Loosen / Tighten' slider post-generation: Loosen adds natural conversational texture (false starts, direct address, more informal language); Tighten makes it crisper and more structured

6.4 Feedback & Learning Loop

- System MUST track three feedback signals per output: Full Accept (user exports as-is), Edit (user modifies the output before exporting), Reject (user requests a full regeneration)
- System MUST capture what was edited in Edit sessions — the delta between what the agent produced and what the user accepted — and use it to update Voice Profile weights
- System MUST offer explicit feedback prompts if user rejects: 'What was wrong with this draft?' with options: Too formal / Too casual / Wrong energy / Wrong structure / Missed the point / Other
- System SHOULD surface a monthly 'Voice Report': 'Your profile has been updated. Here is what changed and what it means for your output.'

6.5 Export & Delivery

- System MUST support export as: plain text (copy to clipboard), formatted PDF with timing markers, Google Doc (via OAuth), Notion page (via integration)
- System SHOULD support direct teleprompter export — formatted output pushed to a linked teleprompter app at the correct line-break cadence for the user's reading pace
- System SHOULD support a 'Content Variants' mode: generate a 60-second version and a written-post version of the same distilled idea from one session

7. The 3-Minute Architecture

Three minutes is not arbitrary — it is the standard length for a strong standalone piece of creator content (YouTube short, Instagram Reel extended, TikTok, standalone vlog segment). Distill treats it as a structural constraint, not a target.

Average speaking pace calibration:

Slow / deliberate	90-110 wpm	270-330 words	Storytellers, reflective creators, emotional content
Conversational	120-140 wpm	360-420 words	Most vloggers, general creators — default calibration
Fast / high energy	150-170 wpm	450-510 words	Comedy, motivation, rapid-fire takes
Variable (user-specific	Calculated	Profile-derived	Distill calculates from user's own content automatically

The 3-minute structure Distill defaults to:

Hook	0:00 - 0:20	~43 words	Grab attention. State or imply the core idea. Create a reason to keep watching.
Bridge	0:20 - 0:35	~33 words	Establish context, credibility, or personal connection to the idea.
Core Point 1	0:35 - 1:10	~76 words	First main idea. Story, example, or evidence. Concrete.
Core Point 2	1:10 - 1:50	~87 words	Second main idea. Should build on or tension-with Point 1.
Core Point 3	1:50 - 2:30	~87 words	Third main idea or pivot. Deepest or most unexpected point.
Closer	2:30 - 3:00	~65 words	Reframe, call to action, or mic-drop line. Must resolve the tension set up in the hook.

This structure is a default, not a rule. The Compression Agent adjusts the structure based on the nature of the content — a story-driven piece may use two long narrative blocks instead of three discrete points, a rapid-fire piece may use 5 short hits. The Voice Profile's structural patterns override the default template.

8. Non-Functional Requirements

Latency	Audio/video transcription: < 60 seconds for up to 20-minute input. Full pipeline (all 4 agents): < 90 seconds from transcription complete to final output. Users see progressive updates — transcription complete, ideas extracted, structure built, voice applied — not a blank screen.
Voice Accuracy	Voice match confidence score must be validated against user feedback. If users consistently rate outputs below 70% match, agent surfaces a recalibration prompt. Target: >80% user satisfaction with voice match after 10+ sessions.
File Support	Audio: mp3, m4a, wav, aac up to 500MB. Video: mp4, mov, webm up to 2GB. Text: plain text, .doc, .docx, .pdf, Google Doc URL. URLs: YouTube, TikTok, Instagram Reels, Vimeo.
Privacy	User's own content (uploaded videos, recordings) is used exclusively for their own Voice Profile. Never used to train shared models. User can delete all ingested content and profile data at any time. Creator reference content is processed but not stored beyond the profile calibration session.
Persistence	Voice Profile stored server-side, synced across devices. Session history (past outputs) retained for 12 months. User can export full history at any time.
Reliability	Pipeline retries each agent step up to 2x on failure before surfacing an error to user. Partial outputs (e.g., transcription complete but compression failed) are preserved so user doesn't lose work.
Accessibility	Full keyboard navigation. Screen reader support for output review. High contrast mode. Font size controls.

9. Technical Architecture

9.1 Stack

- Orchestration: LangGraph (4-agent linear pipeline with conditional branching for low-confidence transcription or low voice-match scenarios)
- LLM: Claude claude-sonnet-4-6 — idea extraction, compression reasoning, voice-calibrated rewriting, 'what got cut' explanation, alternate openers
- Transcription: OpenAI Whisper API (primary), AssemblyAI (fallback for noisy audio — better noise handling)
- URL media extraction: yt-dlp (open source, handles YouTube + TikTok + Vimeo)
- Audio processing: FFmpeg (audio extraction from video, normalization before transcription)
- Backend: Python / FastAPI
- Task Queue: Celery + Redis (async agent pipeline execution, file processing jobs)
- Database: PostgreSQL (Voice Profiles, session history, output drafts, feedback signals), S3 (raw uploaded files, temporary storage during processing)
- Frontend: Next.js / React — input submission, profile setup wizard, output review, diff view, export
- Integrations: Google Docs API (export), Notion API (export), Stripe (billing)
- Infrastructure: AWS ECS Fargate, RDS Postgres, ElastiCache Redis, S3

9.2 Voice Profile Data Model

profile_id	UUID	Primary key
user_id	UUID	Owner
profile_name	String	e.g., 'Main Voice', 'Educational Series'
vocabulary_level	Enum	simple / conversational / technical / mixed
avg_sentence_length	Float	Average words per sentence across own content
sentence_rhythm	Enum	punchy / flowing / mixed

energy_signature	Enum	high / calm / builds / deadpan
opener_pattern	String	Descriptive label + example extracted from own content
closer_pattern	String	Descriptive label + example extracted from own content
pacing_wpm	Integer	Calculated from own content audio
filler_patterns	String[]	List of identified verbal tics and recurring phrases
cultural_markers	String[]	Detected slang, references, humor type
direct_address_rate	Float	Frequency of 'you/we' per 100 words
own_content_count	Integer	Number of own-content items ingested
reference_content_ids	UUID[]	Reference items ingested
reference_weights	JSON	Per-attribute influence of reference content
confidence_score	Float	Overall profile confidence (0-1)
version	Integer	Increments with each profile update
plain_summary	Text	Human-readable profile description
updated_at	DateTime	Last update timestamp

10. Business Model

Free	$0/mo	5 distillations/month, 1 Voice Profile (up to 3 own videos), text input only, basic export (copy/paste)	New creators testing the product
Creator	$19/mo	Unlimited distillations, 1 Voice Profile (unlimited content), all input types (audio, video, URL), all export formats, feedback learning loop	Solo creators posting regularly
Pro	$39/mo	Everything in Creator + 3 Voice Profiles (e.g., different show formats), content variants (60s + written post from same session), teleprompter export, priority processing	Multi-format creators, serious operators
Studio	$79/mo	Everything in Pro + 10 Voice Profiles, team seat (1 collaborator), batch processing (up to 10 inputs queued), API access for custom integrations	Creator teams, production companies, agencies

Pricing anchors to Descript ($24/mo) and Notion AI ($16/mo) — tools creators already pay for. Distill is priced as a premium creative workflow tool, not an AI writing subscription.

11. User Stories

Core Workflow

- As a creator, I want to drop a voice memo I recorded in the car into Distill and get back a 3-minute script in my own voice within 90 seconds so I can film it that evening.
- As a creator, I want to paste a YouTube URL of my favorite creator as a reference so that my output captures their directness without sounding like a copy.
- As a creator, I want to see exactly what got cut from my ramble and why so I can decide if the agent made the right call before I film.
- As a creator, I want two alternate hook options so I don't have to commit to the first one if it doesn't feel right.

Voice Profile

- As a creator, I want to read a plain-language summary of my Voice Profile after setup so I can confirm the agent actually understands how I talk.
- As a creator with two different content series, I want separate Voice Profiles so my educational content sounds different from my vlog content.
- As a new creator with no published content yet, I want to build my Voice Profile from 3 reference creators I admire so I have a starting point while my own style develops.

Feedback & Refinement

- As a creator, I want the agent to learn every time I edit its output so that over time it requires less and less editing.
- As a creator, I want to see my voice match score go up over time as proof the product is actually learning me.

12. Out of Scope — V1

- Video editing — Distill produces scripts and briefs, not edited video. No timeline editor, no caption generation, no B-roll suggestion.
- Live recording — the product does not record audio or video directly. It receives uploaded or linked content.
- Real-time transcription during filming — processing happens offline after submission, not live.
- Music / rap / lyrical content — the 3-minute constraint and structure are calibrated for spoken word/vlog content. Lyrical timing and rhyme scheme require a separate agent architecture.

- Long-form compression (>30 minutes input) — V1 accepts inputs up to 20 minutes. Longer-form documentary or podcast compression is V2.
- Collaboration features — co-creator editing, comment threads, shared Voice Profiles. Studio tier includes one collaborator seat but full team collaboration is V2.

13. Risks & Mitigations

Voice Profile not accurate enough early — first outputs feel generic, users churn before the model improves	High	High	Minimum content threshold before activating profile (3 videos). Low-confidence profile clearly labeled. Free tier limits give users a reason to upload more content. Quick wins: even partial voice calibration is demonstrably better than generic AI output.
Users upload copyrighted reference creator content they don't own	High	Medium	Reference content is processed and discarded — not stored. Platform TOS prohibits copyright infringement. Reference content is used only for structural/tonal pattern extraction, not reproduced.
Transcription quality drops for non-standard accents or heavy background noise	Medium	High	AssemblyAI fallback for noisy audio. Transcription preview step before pipeline runs — user can correct errors. Flag low-confidence transcription clearly.
Output quality plateau — users feel the voice model stops improving after initial calibration	Medium	Medium	Explicit feedback loops (accept/edit/reject) with edit-delta capture. Monthly profile update notifications. 'Recalibrate' option triggers a fresh ingestion session.
yt-dlp breaks due to YouTube policy changes	High (frequent)	Medium	Use Apify or ScraperAPI as fallback for URL extraction. Monitor yt-dlp updates. Users can always upload files directly as a workaround.
Privacy concern — creators don't trust uploading their unpublished content to a third-party platform	Medium	High	Clear privacy policy: own content used only for their profile, never for training shared models. Option for local processing mode in Pro tier (V2). Delete-all feature prominently placed.

14. Open Questions

- Local processing: Should V2 include a local/offline mode where Voice Profile analysis runs on-device rather than server-side? Would significantly help privacy-sensitive power users.
- Teleprompter integrations: Which teleprompter apps (Teleprompter Premium, Parrot Teleprompter, PromptSmart) have open APIs or deep-link support for direct export?
- Reference content rights: Does the Voice Profile's use of reference creator content (for structural pattern extraction) require any licensing disclosure or creator credit?
- Voice match confidence scoring: How is the 87% confidence number actually calculated? Is it based on attribute-by-attribute scoring, user feedback history, or both? The methodology needs to be defined and validated against user perception.
- Content variants scope: In Pro tier, the 60-second variant and written post are derived from the same distillation. Does the written post need its own voice calibration (written voice vs. spoken voice often differ significantly)?
- Onboarding for users with no existing content: The minimum 3-video requirement creates friction for new creators. Is there a bootstrap path using only reference content plus a structured interview (user answers 5 questions on camera) as a substitute?

The Compounding Advantage

Distill gets better with every session. The first output is good. The tenth output is yours. That compounding effect — where the product becomes more valuable the more you use it — is the core of the retention model. Every creator who uses Distill for 30 days has a Voice Profile that took 30 days to build. That profile doesn't exist anywhere else. That is the moat.

This document is a working draft. Technical specifications are directional and subject to revision during engineering scoping.